

hpc_hw2

March 24, 2017

1 Christopher Prince [cmp670@nyu.edu]

1.1 MATH.2012 HW 2

1.1.1 Preliminaries

```
In [1]: import pandas as pd
import numpy as np
import os
```

```
In [2]: import pylab as pl
%matplotlib inline
```

1.2 Source files and reproducibility

The program has been packaged with a make file and a bash script to run the experiment on a given machine. The process is:

- Clone the github repo cmprince/hpc17 (files are in the hw2 directory).
- Run make.
- Run runlaplace.sh (if necessary change file permissions). This runs both Gauss-Seidel and Jacobi iterations for $N \in \{50, 100, 1000\}$, repeating 5 times, and writes the results of each run to a file.

1.3 Architectures

The tests were run on 4 different machines: my personal laptop, a server at NYU-CUSP, and two servers at CIMS. The specifications for the CUSP server was obtained from datahub.cusp.nyu.edu, and for the CIMS server from cims.nyu.edu. Additionally, on brawler, versions of the codes without parallizing the residual calculations were run seperately to see the effect on performance.

Machine	Owner	Processor	No. Processors	No. Cores/Proc.	Freq. (GHz)	Memory (GB)	O
brawler	CIMS	AMD Opteron	2	2	2.6	8	R
compute	CUSP	Xeon E5-4640	4	8	2.40	1024	O
crunchy4	CIMS	AMD Opteron 6136	4	8	2.4	128	R
laptop	Me	Xeon E3-1505M v5	1	4	2.80	64	U

1.4 Results and observations

1.4.1 Timing results

The results files are space-delimited with the following columns:

```
In [3]: names=['N', 'run', 'usec']
```

Note timings are reported in μsec . Uncomment the cell below to retrieve the sample data used in this notebook:

```
In [4]: #!/wget -r -l1 -np -nd -R index.html?*,html -e robots=off https://www.cims.nyu.edu/~cmp67
```

Read the files in as dataframes and merge the timing columns, starting with Jacobi:

```
In [5]: gb = pd.read_csv("jacobi2D-omp_laptop", header=None, delimiter=" ", names=names).drop("u
```

```
In [6]: for f in sorted(os.listdir('.')):
        if f[:13]=="jacobi2D-omp_":
            g = pd.read_csv(f, header=None, delimiter=" ", names=names[:-1] + [f])
            gb = gb.merge(g, on=names[:-1])
```

```
In [7]: gb
```

```
Out[7]:
```

	N	run	jacobi2D-omp_brawler	jacobi2D-omp_brawler_no_l2norm_parallel	\
0	50	1	64098	122950	
1	50	2	62954	324677	
2	50	3	62743	122595	
3	50	4	64429	122490	
4	50	5	62655	122655	
5	100	1	225179	492691	
6	100	2	224954	685548	
7	100	3	230868	1062484	
8	100	4	230655	596687	
9	100	5	224027	579145	
10	1000	1	21845399	47086142	
11	1000	2	22503545	47557294	
12	1000	3	22449011	47027564	
13	1000	4	21808371	46989342	
14	1000	5	21793035	47565443	
			jacobi2D-omp_crunchy4	jacobi2D-omp_laptop	jacobi2D-omp_mauler
0			3869467	79251	64626
1			3903125	25442	62799
2			3926278	24723	64779
3			3864763	21033	64820
4			3871138	20687	63231
5			3942070	73839	225253
6			3914154	72827	225845
7			405486	76498	225949

8	3895513	72453	230387
9	3924761	78413	225295
10	6944077	7032157	22376295
11	6916533	6291400	23219363
12	7054551	6769999	22188391
13	5993197	6500057	21984691
14	7056586	6462039	21998788

There are a few runs where the timing is much greater than the others; instead of taking a mean value over the runs I calculate the median run time here.

```
In [8]: g = gb.groupby(names[0])
        g.agg(np.median).drop('run', axis=1)
```

```
Out [8]:      jacobi2D-omp_brawler  jacobi2D-omp_brawler_no_l2norm_parallel \
N
50          62954          122655
100         225179          596687
1000        21845399          47086142

      jacobi2D-omp_crunchy4  jacobi2D-omp_laptop  jacobi2D-omp_mauler
N
50          3871138          24723          64626
100         3914154          73839          225845
1000        6944077          6500057          22188391
```

Repeating for Gauss-Seidel:

```
In [9]: gb = pd.read_csv("gs2D-omp_laptop", header=None, delimiter=" ", names=names).drop("usec"
```

```
In [10]: for f in sorted(os.listdir('.')):
          if f[:9]=="gs2D-omp_":
              g = pd.read_csv(f, header=None, delimiter=" ", names=names[:-1] + [f])
              gb = gb.merge(g, on=names[:-1])
```

```
In [11]: gb
```

```
Out [11]:      N  run  gs2D-omp_brawler  gs2D-omp_crunchy4  gs2D-omp_laptop \
0    50    1          124188          3893753          85588
1    50    2          125482          3838401          85986
2    50    3          124588          3890386          80742
3    50    4          124582          3876093          84353
4    50    5          125483          3887672          83798
5   100    1          483850          3859820          309043
6   100    2          841047          3898532          310122
7   100    3          483088          2820291          311661
8   100    4          483126          3886792          314079
9   100    5          489542          3962736          313914
10  1000    1          46242520          32425595          25087905
```

11	1000	2	47183560	32286412	25479990
12	1000	3	46306250	32326000	25177850
13	1000	4	46119539	32287583	23282460
14	1000	5	46696857	32292971	23516537

	gs2D-omp_mauler	gs2D-omp_mauler_no_l2norm_parallel
0	126195	134271
1	211472	126126
2	124792	179788
3	124500	125224
4	125350	124929
5	484326	487323
6	1353349	490656
7	497028	497853
8	483626	484107
9	580930	486020
10	46590516	47566617
11	46820416	47573191
12	46644781	46385122
13	46665327	46865096
14	47380288	46662648

```
In [12]: g = gb.groupby(names[0])
g.agg(np.median).drop('run', axis=1)
```

```
Out[12]:
```

	gs2D-omp_brawler	gs2D-omp_crunchy4	gs2D-omp_laptop	gs2D-omp_mauler	\
N					
50	124588	3887672	84353	125350	
100	483850	3886792	311661	497028	
1000	46306250	32292971	25087905	46665327	

	gs2D-omp_mauler_no_l2norm_parallel
N	
50	126126
100	487323
1000	46865096